

Oberry is losing members quietly.

~19%

of members go silent for 60+ days, unnoticed

2,260

members already flagged high-risk today

฿5.4M/yr


in spend riding on reaching them in time

No cancellation, no complaint — a member just stops coming back, and nothing in a typical POS report flags it until they're already gone.


One pipeline: data → segments → risk score → action





Two years of transaction history becomes five member segments, a trained risk score for every member, and a specific recommended action — not a hunch.



Five kinds of member, not one undifferentiated list




Does it work? Yes — with honest, stated limits



At the chosen threshold, the model catches 57% of members who actually leave, at 70% precision on every alarm it raises — no model catches everyone, and this one doesn't pretend to.

The strongest warning isn't spend — it's silence



- A widening gap between visits predicts departure weeks before the last one.
- Members who never redeem a point churn noticeably more often.
- Recency and visit-gap trend dominate — spend alone barely registers.

A specific action for every flagged member

Segment	Risk	Recommended action
● Champions	low	VIP perks, early access to new drinks, referral ask
● Loyal	low–med	Tier-up nudge, personalized bundle
● At-risk regulars	high	15%-off win-back coupon + “we miss you” LINE mission
● Hibernating	very high	Bounce-back free drink + one-question why-survey
● One-and-done	high	Onboarding mission, second-visit nudge

Worked example — member M016486

Bronze tier, Champions segment, \$19,563/year in spend. Typical gap between visits is now 18.2x longer than their own usual pattern (+27 points of risk). Risk: 99.6%. Action: VIP perks, early access to new drinks, referral ask.

What it's worth, and what's next

20,000

total members

→ 14,337

scored by the model

→ 2,260

flagged high-risk

→ £5.4M/yr

in spend on the line

To productionize, this needs:

- Oberry's real transaction data, not synthetic
- A scheduled batch-scoring run, not a one-off
- Instrumented campaigns with measured lift
- Ongoing monitoring for model drift

This is a working prototype, validated on synthetic data end to end — the architecture is real and the pipeline runs today. What's unproven is whether these exact patterns hold in Oberry's real transaction history, and that's the next step, not a leap of faith.